

# Week 3 Questions (220962018)

January 25, 2025

## 0.1 Week 3 Lab Exercise

```
[67]: import torch
      from matplotlib import pyplot as plt
      from torch import nn
      import torch.optim as optim
      from torch.utils.data import Dataset, DataLoader
      import numpy as np
      device = "cuda" if torch.cuda.is_available() else "cpu"
      torch.cuda.device_count() , device
```

[67]: (1, 'cuda')

## 0.2 Lab Exercise

**0.2.1 1) For the following training data, build a linear regression model. Assume  $w$  and  $b$  are initialized with 1 and learning parameter is set to 0.001.**

$x = \text{torch.tensor}([12.4, 14.3, 14.5, 14.9, 16.1, 16.9, 16.5, 15.4, 17.0, 17.9, 18.8, 20.3, 22.4, 19.4, 15.5, 16.7, 17.3, 18.4, 19.2, 17.4, 19.5, 19.7, 21.2])$

$y = \text{torch.tensor}([11.2, 12.5, 12.7, 13.1, 14.1, 14.8, 14.4, 13.4, 14.9, 15.6, 16.4, 17.7, 19.6, 16.9, 14.0, 14.6, 15.1, 16.1, 16.8, 15.2, 17.0, 17.2, 18.6])$

Assume learning rate = 0.001. Plot the graph of epoch in x axis and loss in y axis.

```
[68]: x=torch.tensor([12.4, 14.3, 14.5, 14.9, 16.1, 16.9, 16.5, 15.4, 17.0, 17.9, 18.
      ↪8, 20.3, 22.4,
      19.4, 15.5, 16.7, 17.3, 18.4, 19.2, 17.4, 19.5, 19.7, 21.2])
      y=torch.tensor([11.2, 12.5, 12.7, 13.1, 14.1, 14.8, 14.4, 13.4, 14.9, 15.6, 16.
      ↪4, 17.7, 19.6,
      16.9, 14.0, 14.6, 15.1, 16.1, 16.8, 15.2, 17.0, 17.2, 18.6])
```

```
[69]: def mse(y, y_pred):
      return torch.mean((y - y_pred)**2)

      lr = 0.001
      epochs = 20
```

```

w, b = torch.tensor(1., requires_grad=True), torch.tensor(1.,
↪requires_grad=True)
losses = []

for i in range(epochs):
    y_pred = w * x + b

    loss = mse(y, y_pred)

    loss.backward()

    with torch.no_grad():
        w -= lr * w.grad
        b -= lr * b.grad
        losses.append(loss.detach())
        print(loss)
    w.grad.zero_()
    b.grad.zero_()


plt.plot(list(range(1,epochs+1)), losses)
plt.xlabel('No. of Epochs')
plt.ylabel('Loss')
plt.show()

```

```

tensor(10.1539, grad_fn=<MeanBackward0>)
tensor(1.4567, grad_fn=<MeanBackward0>)
tensor(0.2217, grad_fn=<MeanBackward0>)
tensor(0.0464, grad_fn=<MeanBackward0>)
tensor(0.0215, grad_fn=<MeanBackward0>)
tensor(0.0179, grad_fn=<MeanBackward0>)
tensor(0.0174, grad_fn=<MeanBackward0>)
tensor(0.0174, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)
tensor(0.0173, grad_fn=<MeanBackward0>)

```



**0.2.2 2)** Find the value of  $w.grad$ ,  $b.grad$  using analytical solution for the given linear regression problem. Initial value of  $w = b = 1$ . Learning parameter is set to 0.001. Implement the same and verify the values of  $w.grad$ ,  $b.grad$  and updated parameter values for two epochs. Consider the difference between predicted and target values of  $y$  is defined as  $(y_p - y)$ .

```
[70]: lr = 0.001
epochs = 2
w, b = torch.tensor(1., requires_grad=True), torch.tensor(1., requires_grad=True)

x = torch.tensor([2, 4])
y = torch.tensor([20, 40])

for i in range(epochs):
    print(w, b)
    y_pred = w * x + b

    loss = mse(y, y_pred)

    loss.backward()
```

```

print(f"Epoch {i+1}: w_grad={w.grad.item()}, b_grad={b.grad.item()}")

with torch.no_grad():
    w -= lr * w.grad
    b -= lr * b.grad

w.grad.zero_()
b.grad.zero_()

print(f"Epoch {i+1}: Loss={loss.item()}\n")

```

w, b

```
tensor(1., requires_grad=True) tensor(1., requires_grad=True)
```

```
Epoch 1: w_grad=-174.0, b_grad=-52.0
```

```
Epoch 1: Loss=757.0
```

```
tensor(1.1740, requires_grad=True) tensor(1.0520, requires_grad=True)
```

```
Epoch 2: w_grad=-170.20799255371094, b_grad=-50.85199737548828
```

```
Epoch 2: Loss=724.3797607421875
```

```
[70]: (tensor(1.3442, requires_grad=True), tensor(1.1029, requires_grad=True))
```

**0.2.3 3) Revise the linear regression model by defining a user defined class titled RegressionModel with two parameters w and b as its member variables. Define a constructor to initialize w and b with value 1. Define four member functions namely forward(x) to implement  $w x + b$ , update() to update w and b values, reset\_grad() to reset parameters to zero, criterion(y, yp) to implement MSE Loss given the predicted y value yp and the target label y. Define an object of this class named model and invoke all the methods. Plot the graph of epoch vs loss by varying epoch to 100 iterations.**

```
x = torch.tensor([5.0, 7.0, 12.0, 16.0, 20.0])
```

```
y = torch.tensor([40.0, 120.0, 180.0, 210.0, 240.0])
```

```
learning_rate = torch.tensor(0.001)
```

```

[71]: class RegressionModel:
    def __init__(self, w, b, lr=0.001):
        self.w = torch.tensor(w, requires_grad=True)
        self.b = torch.tensor(b, requires_grad=True)
        self.lr = lr
        self.y = None
        self.loss = None
        self.losses = []

```

```

def forward(self, x):
    self.y = self.w * x + self.b
    return self.y

def update(self):
    self.loss.backward()
    with torch.no_grad():
        self.w -= self.lr * self.w.grad
        self.b -= self.lr * self.b.grad

def reset_grad(self):
    self.w.grad.zero_()
    self.b.grad.zero_()

def criterion(self, y, yp):
    self.loss = torch.mean((y - y_pred)**2)
    self.losses.append(self.loss.detach())
    return self.loss

```

```

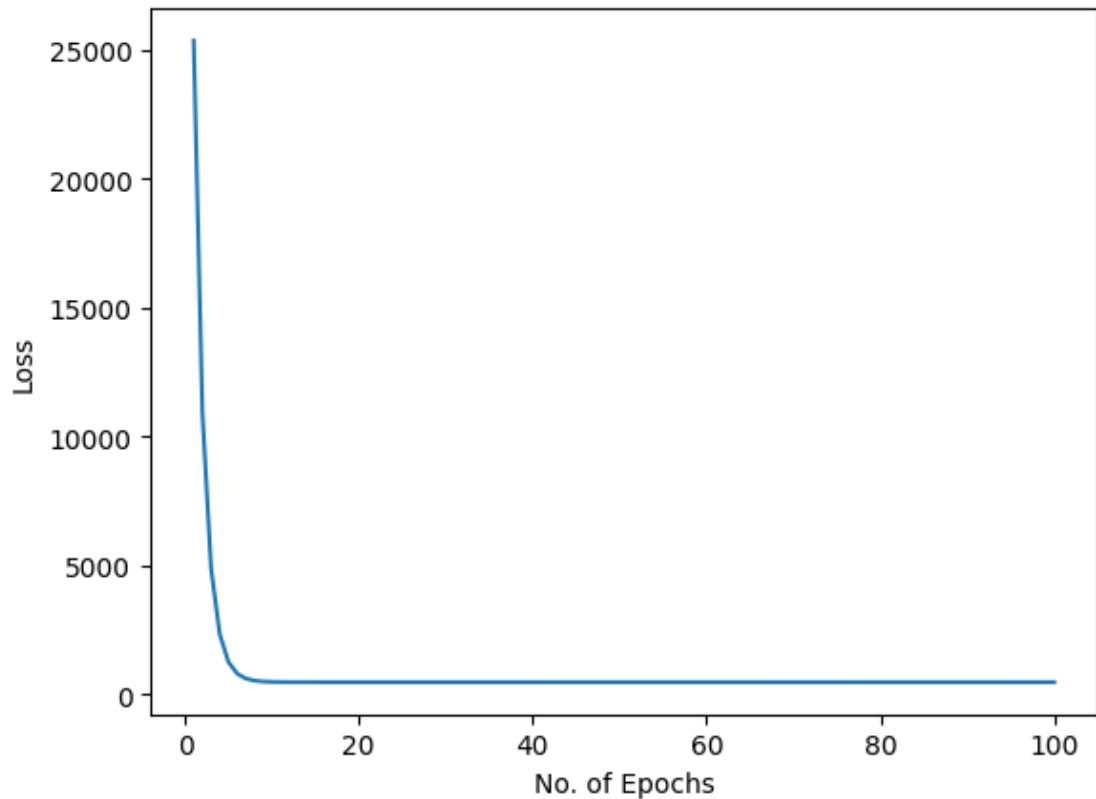
[72]: model = RegressionModel(1., 1.)
epochs = 100

x = torch.tensor([5.0, 7.0, 12.0, 16.0, 20.0])
y = torch.tensor([40.0, 120.0, 180.0, 210.0, 240.0])

for i in range(epochs):
    y_pred = model.forward(x)
    loss = model.criterion(y, y_pred)
    model.update()
    model.reset_grad()

plt.plot(np.arange(1, epochs+1), model.losses)
plt.xlabel('No. of Epochs')
plt.ylabel('Loss')
plt.show()

```



**0.2.4 4) Convert your program written in Qn 3 to extend `nn.module` in your model. Also override the necessary methods to fit the regression line. Illustrate the use of `Dataset` and `DataLoader` from `torch.utils.data` in your implementation. Use the SGD Optimizer `torch.optim.SGD()`**

```
[73]: def criterion(yp, y):
        return torch.mean((y - y_pred)**2)

x = torch.tensor([5.0, 7.0, 12.0, 16.0, 20.0])
y = torch.tensor([40.0, 120.0, 180.0, 210.0, 240.0])

w = 1.
b = 1.

# Define the RegressionModel class inheriting from nn.Module
class RegressionModel(nn.Module):
    def __init__(self, w, b):
        super(RegressionModel, self).__init__()
        self.w = nn.Parameter(torch.tensor(w, requires_grad=True, dtype=torch.
        ↪float))
```

```

        self.b = nn.Parameter(torch.tensor(b, requires_grad=True, dtype=torch.
↪float))

    def forward(self, x):
        return self.w * x + self.b

# Custom Dataset to handle input and output data
class LinearDataset(Dataset):
    def __init__(self, x_data, y_data):
        self.x_data = x_data
        self.y_data = y_data

    def __len__(self):
        return len(self.x_data)

    def __getitem__(self, idx):
        return self.x_data[idx], self.y_data[idx]

# Create the dataset and data loader
dataset = LinearDataset(x, y)
dataloader = DataLoader(dataset, batch_size=16, shuffle=False)

# Initialize the model, loss function, and optimizer
model = RegressionModel(w, b)
#criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=1e-4)

# Training loop
num_epochs = 100
for epoch in range(num_epochs):
    for batch_x, batch_y in dataloader:
        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        y_pred = model(batch_x)

        # Compute loss
        loss = criterion(y_pred, batch_y)

        # Backward pass
        loss.backward()

        # Update weights
        optimizer.step()

    if (epoch + 1) % 10 == 0:

```

```
print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}')
```

```
Epoch [10/100], Loss: 13552.6689
Epoch [20/100], Loss: 6875.4150
Epoch [30/100], Loss: 3609.3821
Epoch [40/100], Loss: 2011.8713
Epoch [50/100], Loss: 1230.4796
Epoch [60/100], Loss: 848.2745
Epoch [70/100], Loss: 661.3213
Epoch [80/100], Loss: 569.8713
Epoch [90/100], Loss: 525.1349
Epoch [100/100], Loss: 503.2474
```

**0.2.5 5) Use PyTorch's `nn.Linear()` in your implementation to perform linear regression for the data provided in Qn. 1. Also plot the graph.**

```
[74]: x = torch.tensor( [12.4, 14.3, 14.5, 14.9, 16.1, 16.9, 16.5, 15.4, 17.0, 17.9,
    ↪18.8, 20.3, 22.4,
    19.4, 15.5, 16.7, 17.3, 18.4, 19.2, 17.4, 19.5, 19.7, 21.2])

y = torch.tensor( [11.2, 12.5, 12.7, 13.1, 14.1, 14.8, 14.4, 13.4, 14.9, 15.6,
    ↪16.4, 17.7, 19.6,
    16.9, 14.0, 14.6, 15.1, 16.1, 16.8, 15.2, 17.0, 17.2, 18.6])

losses = []

# Create the dataset and data loader
dataset = LinearDataset(x, y)
dataloader = DataLoader(dataset, batch_size=1, shuffle=True)

# Initialize the model, loss function, and optimizer
model = nn.Linear(1,1)
# criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.0001)

# Training loop
num_epochs = 100
for epoch in range(num_epochs):
    for batch_x, batch_y in dataloader:
        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        y_pred = model(batch_x)

        # Compute loss
        loss = criterion(y_pred, batch_y)
```



```

        # Backward pass
        loss.backward()

        # Update weights
        optimizer.step()

    losses.append(loss.item())

    if (epoch + 1) % 10 == 0:
        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}')


plt.plot(np.arange(1, num_epochs + 1), losses)
plt.xlabel('No. of Epochs')
plt.ylabel('Loss')
plt.show()

```

```

Epoch [10/100], Loss: 0.0006
Epoch [20/100], Loss: 0.0000
Epoch [30/100], Loss: 0.0025
Epoch [40/100], Loss: 0.0008
Epoch [50/100], Loss: 0.0007
Epoch [60/100], Loss: 0.3828
Epoch [70/100], Loss: 0.0341
Epoch [80/100], Loss: 0.0024
Epoch [90/100], Loss: 0.0002
Epoch [100/100], Loss: 0.0321

```



**0.2.6 6) Implement multiple linear regression for the data provided below**

Subject	X1	X2	Y
1	3	8	-3.7
2	4	5	3.5
3	5	7	2.5
4	6	3	11.5
5	2	1	5.7

Verify your answer for the data point X1=3, X2=2.

```
[75]: x1 = torch.tensor([3, 4, 5, 6, 2])
x2 = torch.tensor([8, 5, 7, 3, 1])
x = torch.stack((x1, x2), dim=1)
y = torch.tensor([-3.7, 3.5, 2.5, 11.5, 5.7])

losses = []

epochs = 20
w, b = [1., 1.], 1.
```

```

# Create the dataset and data loader
dataset = LinearDataset(x, y)
dataloader = DataLoader(dataset, batch_size=1, shuffle=False)

# Initialize the model, loss function, and optimizer
model = RegressionModel(w, b)
# criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=1e-2)

# Training loop
num_epochs = 100
for epoch in range(num_epochs):
    for batch_x, batch_y in dataloader:
        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        y_pred = model(batch_x)

        # Compute loss
        loss = criterion(y_pred, batch_y)

        # Backward pass
        loss.backward()

        # Update weights
        optimizer.step()

    losses.append(loss.detach())
    if (epoch + 1) % 1 == 0:
        print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}')

plt.plot(np.arange(1, num_epochs+1), losses)
plt.xlabel('No. of Epochs')
plt.ylabel('Loss')
plt.show()

```


```

Epoch [1/100], Loss: 12.3834
Epoch [2/100], Loss: 11.9640
Epoch [3/100], Loss: 11.5118
Epoch [4/100], Loss: 11.1089
Epoch [5/100], Loss: 10.7362
Epoch [6/100], Loss: 10.3819
Epoch [7/100], Loss: 10.0413
Epoch [8/100], Loss: 9.7126
Epoch [9/100], Loss: 9.3948
Epoch [10/100], Loss: 9.0875

```

Epoch [11/100], Loss: 8.7903  
Epoch [12/100], Loss: 8.5028  
Epoch [13/100], Loss: 8.2247  
Epoch [14/100], Loss: 7.9558  
Epoch [15/100], Loss: 7.6957  
Epoch [16/100], Loss: 7.4440  
Epoch [17/100], Loss: 7.2007  
Epoch [18/100], Loss: 6.9653  
Epoch [19/100], Loss: 6.7377  
Epoch [20/100], Loss: 6.5175  
Epoch [21/100], Loss: 6.3045  
Epoch [22/100], Loss: 6.0985  
Epoch [23/100], Loss: 5.8993  
Epoch [24/100], Loss: 5.7067  
Epoch [25/100], Loss: 5.5203  
Epoch [26/100], Loss: 5.3401  
Epoch [27/100], Loss: 5.1658  
Epoch [28/100], Loss: 4.9972  
Epoch [29/100], Loss: 4.8342  
Epoch [30/100], Loss: 4.6766  
Epoch [31/100], Loss: 4.5241  
Epoch [32/100], Loss: 4.3766  
Epoch [33/100], Loss: 4.2340  
Epoch [34/100], Loss: 4.0961  
Epoch [35/100], Loss: 3.9627  
Epoch [36/100], Loss: 3.8338  
Epoch [37/100], Loss: 3.7090  
Epoch [38/100], Loss: 3.5884  
Epoch [39/100], Loss: 3.4718  
Epoch [40/100], Loss: 3.3590  
Epoch [41/100], Loss: 3.2500  
Epoch [42/100], Loss: 3.1445  
Epoch [43/100], Loss: 3.0425  
Epoch [44/100], Loss: 2.9439  
Epoch [45/100], Loss: 2.8485  
Epoch [46/100], Loss: 2.7563  
Epoch [47/100], Loss: 2.6672  
Epoch [48/100], Loss: 2.5810  
Epoch [49/100], Loss: 2.4976  
Epoch [50/100], Loss: 2.4170  
Epoch [51/100], Loss: 2.3391  
Epoch [52/100], Loss: 2.2638  
Epoch [53/100], Loss: 2.1909  
Epoch [54/100], Loss: 2.1205  
Epoch [55/100], Loss: 2.0524  
Epoch [56/100], Loss: 1.9866  
Epoch [57/100], Loss: 1.9229  
Epoch [58/100], Loss: 1.8614

Epoch [59/100], Loss: 1.8019  
Epoch [60/100], Loss: 1.7444  
Epoch [61/100], Loss: 1.6888  
Epoch [62/100], Loss: 1.6350  
Epoch [63/100], Loss: 1.5831  
Epoch [64/100], Loss: 1.5329  
Epoch [65/100], Loss: 1.4843  
Epoch [66/100], Loss: 1.4374  
Epoch [67/100], Loss: 1.3920  
Epoch [68/100], Loss: 1.3481  
Epoch [69/100], Loss: 1.3057  
Epoch [70/100], Loss: 1.2648  
Epoch [71/100], Loss: 1.2252  
Epoch [72/100], Loss: 1.1869  
Epoch [73/100], Loss: 1.1499  
Epoch [74/100], Loss: 1.1141  
Epoch [75/100], Loss: 1.0796  
Epoch [76/100], Loss: 1.0462  
Epoch [77/100], Loss: 1.0139  
Epoch [78/100], Loss: 0.9827  
Epoch [79/100], Loss: 0.9525  
Epoch [80/100], Loss: 0.9234  
Epoch [81/100], Loss: 0.8952  
Epoch [82/100], Loss: 0.8680  
Epoch [83/100], Loss: 0.8417  
Epoch [84/100], Loss: 0.8163  
Epoch [85/100], Loss: 0.7918  
Epoch [86/100], Loss: 0.7681  
Epoch [87/100], Loss: 0.7451  
Epoch [88/100], Loss: 0.7230  
Epoch [89/100], Loss: 0.7016  
Epoch [90/100], Loss: 0.6809  
Epoch [91/100], Loss: 0.6610  
Epoch [92/100], Loss: 0.6417  
Epoch [93/100], Loss: 0.6230  
Epoch [94/100], Loss: 0.6050  
Epoch [95/100], Loss: 0.5876  
Epoch [96/100], Loss: 0.5708  
Epoch [97/100], Loss: 0.5546  
Epoch [98/100], Loss: 0.5389  
Epoch [99/100], Loss: 0.5238  
Epoch [100/100], Loss: 0.5092



### 0.2.7 7) Implement logistic regression

x = [1, 5, 10, 10, 25, 50, 70, 75, 100,]

y = [0, 0, 0, 0, 0, 1, 1, 1, 1]

```
[76]: class LogisticModel(nn.Module):
    def __init__(self, w, b):
        super(LogisticModel, self).__init__()
        self.w = nn.Parameter(torch.tensor(w, requires_grad=True, dtype=torch.
↪float))
        self.b = nn.Parameter(torch.tensor(b, requires_grad=True, dtype=torch.
↪float))

    def forward(self, x):
        z = self.w * x + self.b
        #print("stuff:", self.w, x, self.b)
        return torch.sigmoid(z) # 1 / (1 + torch.exp(-z))

x = torch.tensor([1, 5, 10, 10, 25, 50, 70, 75, 100], dtype=torch.float)
y = torch.tensor([0, 0, 0, 0, 0, 1, 1, 1, 1], dtype=torch.float)
```

```

losses = []

epochs = 20
w, b = 1., 1.

# Create the dataset and data loader
dataset = LinearDataset(x, y)
dataloader = DataLoader(dataset, batch_size=1, shuffle=False)

# Initialize the model, loss function, and optimizer
model = LogisticModel(w, b)
criterion = nn.BCELoss()
optimizer = optim.SGD(model.parameters(), lr=0.005)

# Training loop
num_epochs = 200
for epoch in range(num_epochs):
    epoch_loss = []
    for batch_x, batch_y in dataloader:
        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        y_pred = model(batch_x)

        # Compute loss
        loss = criterion(y_pred, batch_y)

        # Backward pass
        loss.backward()
        epoch_loss.append(loss.detach())

        # Update weights
        optimizer.step()

    losses.append(sum(epoch_loss) / len(epoch_loss))
    if (epoch+1) % 10 == 0:
        print(f"Loss for Epoch {epoch+1}/{num_epochs}: {losses[-1]}")

plt.plot(np.arange(1, num_epochs+1), losses)
plt.xlabel('No. of Epochs')
plt.ylabel('Loss')
plt.show()

```

```

Loss for Epoch 10/200: 1.3185741901397705
Loss for Epoch 20/200: 1.2531628608703613
Loss for Epoch 30/200: 1.1903891563415527

```

Loss for Epoch 40/200: 1.1303671598434448  
Loss for Epoch 50/200: 1.0731796026229858  
Loss for Epoch 60/200: 1.0188769102096558  
Loss for Epoch 70/200: 0.9674785137176514  
Loss for Epoch 80/200: 0.9189727306365967  
Loss for Epoch 90/200: 0.8733193278312683  
Loss for Epoch 100/200: 0.8304525017738342  
Loss for Epoch 110/200: 0.7902848124504089  
Loss for Epoch 120/200: 0.7527111768722534  
Loss for Epoch 130/200: 0.7176131010055542  
Loss for Epoch 140/200: 0.6848623752593994  
Loss for Epoch 150/200: 0.6543242931365967  
Loss for Epoch 160/200: 0.6258624792098999  
Loss for Epoch 170/200: 0.5993388295173645  
Loss for Epoch 180/200: 0.5746187567710876  
Loss for Epoch 190/200: 0.5515710711479187  
Loss for Epoch 200/200: 0.5300692915916443

